

Introduction to the use of weights for sampling/selection bias

Lorenzo Fabbri

Xavier Basagana

Martine Vrijheid

2024-03-07

Table of contents

1	Selection bias	1
1.1	A simple example	1
1.2	Some strategies	3
2	Estimating the balancing weights	4
2.1	Step-by-step procedure	4
2.2	The SelectionWeights R package	8
3	Suggestions	8
	References	8

1 Selection bias

From the subjects belonging to the HELIX sub-cohort, only a fraction took part to the follow-up (as part of the ATHLETE project). That is, among the eligible subjects, some are going to be excluded from our analyses since they have e.g., no outcome. Censoring from the analysis those with missing values will most likely introduce **selection bias**.

1.1 A simple example

Suppose exposure **A** is a measure of socio-economic position (SEP), and outcome **Y** is diagnosis of ADHD. The censoring variable **C** is a collider on this pathway, and **L** is a confounder. To estimate the effect of **A** on **Y** we should NOT adjust for **C** (Figure 1).

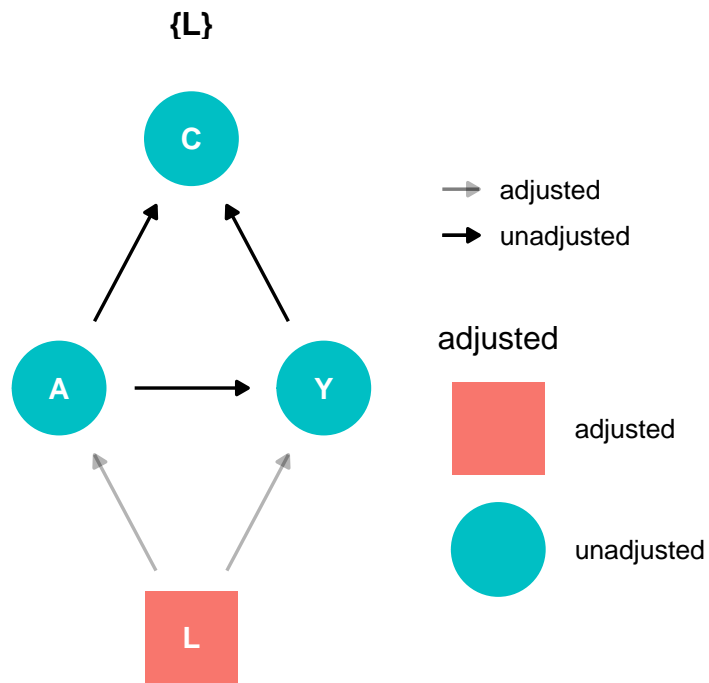


Figure 1: Effect estimation with selection bias

The problem is that in the follow-up, we are implicitly adjusting for \mathbf{C} as well, thus opening another path. Thus, censoring due to loss to follow-up can introduce selection bias. Generally, we are interested then in estimating an effect if nobody had been censored. If \mathbf{A} is binary, we are interested in:

$$E[Y^{a=1, c=0}] - E[Y^{a=0, c=0}],$$

which is the joint effect of \mathbf{A} and \mathbf{C} .

1.2 Some strategies

There are multiple strategies available to reduce the effect of this selection bias on the association between exposure and outcome. In what follows, we assume that the identifiability conditions for the joint treatment (A, C) conditional on L hold.

One such strategy consists in computing **inverse-probability (IP) weights**. In this case, we are interested in estimating the parameters of the marginal structural model (MSM) $E[Y^{a, c=0}] = \beta_0 + \beta_1 a$, with \mathbf{a} being the exposure and \mathbf{c} a indicator variable for censoring. Then, the IP weights are defined as follows:

$$W^{A, C} = W^A \times W^C, \tag{1}$$

with $W^C = 1/Pr[C = 0|L, A]$ ($W^C = 0$ for the censored), and W^A being the weights to adjust for confounding ($f(A|L)$). Alternatively, we can compute *stabilized* IP weights:

$$SW^C = Pr[C = 0|A] / Pr[C = 0|L, A].$$

With IP weights, we are effectively modeling the censoring mechanism given the exposure and a set of covariates (e.g., sex and age). Another strategy consists in directly modeling the outcome (standardization). In this case, instead of using weights to adjust for confounding and selection bias, we are directly including these variables in the outcome model:

$$E[Y|A = a, C = 0, L = l].$$

Another strategy consists in modeling both the censoring mechanism and the outcome, and *combine* the results together. In this case, we would fit a weighted outcome model, where the weights are as in Equation 1. The same set of covariates used to estimate W^C should be incorporated also in the outcome model.

2 Estimating the balancing weights

In what follows, we will focus on how to estimate these balancing weights, and how to use them when fitting the outcome model. Since weight estimation depends on question-specific confounders and eventually exposures, it makes sense for the researcher to estimate them directly, rather than providing them for *general purposes*.

2.1 Step-by-step procedure

We will make use of the `WeightIt` R package to estimate the IP weights ([Greifer 2024b](#)). This package offers a multitude of methods to estimate the balancing weights, from a traditional generalized linear model to data-adaptive methods that do not make use of any parametric specification. The researcher is suggested to try different methods and assess the resulting balance with strategies outlined below.

We will make use of the Lalonde dataset, and we will estimate the effect of `treat` on the continuous outcome `re78`. We will assume that `age`, `educ`, `nodegree`, and `race` are sufficient to remove any confounding bias. We will further create a new (random) variable, `c`, to indicate whether each subject was either censored ($c = 1$) or not ($c = 0$), based on `educ` and `age`. This is a very simplistic scenario to illustrate the basic steps.

```
type_exposure <- "binary"

dat <- cobalt::lalonge |>
  tibble::as_tibble() |>
  tidylog::mutate(
    censoring_mech = age * educ
  )
dat$c <- with(dat, rbinom(
  n = nrow(dat),
  size = 1,
  prob = (censoring_mech - min(censoring_mech)) /
    (max(censoring_mech) - min(censoring_mech))
))
```

We can now proceed to estimate the balancing weights for selection bias, assuming that the censoring mechanism solely involves the variables `educ`, `age`, and `race`. We will use the `energy` method, which is not based on any parametric assumption. Some of the estimated weights are considered *extreme* (you can check this with the `summary` function applied to the returning object, which provides other useful information like the effective sample size before and after weighting). We therefore trim/winsorize them.

```
est_sel_weights <- WeightIt::weightit(
  formula = c ~ educ + age + race,
  data = dat,
  method = "energy"
)
summary(est_sel_weights)
```

Summary of weights

- Weight ranges:

	Min		Max
treated	0.0943	-----	5.4791
control	0.0593	-----	5.1802

- Units with the 5 most extreme weights by group:

	481	538	596	602	463
treated	4.0516	4.0593	4.5411	5.1302	5.4791
	142	310	305	315	549
control	4.5969	4.7439	4.9779	5.1802	5.1802

- Weight statistics:

	Coef of Var	MAD	Entropy	# Zeros
treated	0.882	0.642	0.329	0
control	0.906	0.655	0.336	0

- Effective Sample Sizes:

	Control	Treated
Unweighted	368.	246.
Weighted	202.39	138.65

```
est_sel_weights_trim <- WeightIt::trim(
  est_sel_weights,
  at = 0.7,
  lower = TRUE
)
summary(est_sel_weights_trim)
```

Summary of weights

- Weight ranges:

	Min	Max
treated	0.315	1.2366
control	0.315	1.2366

- Units with the 5 most extreme weights by group:

	52	36	12	11	4
treated	1.2366	1.2366	1.2366	1.2366	1.2366
	19	17	15	14	9
control	1.2366	1.2366	1.2366	1.2366	1.2366

- Weight statistics:

	Coef of Var	MAD	Entropy	# Zeros
treated	0.491	0.454	0.129	0
control	0.524	0.493	0.146	0

- Effective Sample Sizes:

	Control	Treated
Unweighted	368.	246.
Weighted	288.91	198.39

We can visually check whether the estimated weights are effectively capable of reducing the effect of the chosen covariates on the censoring mechanism using a so-called *Love plot*. Ideally, we would like to see that the adjusted values are as close as possible to zero. Deviations from this should be solved for instance by choosing another fitting method and/or by modifying the trimming value.

```
cobalt::love.plot(  
  est_sel_weights_trim,  
  stats = ifelse(  
    type_exposure == "continuous",  
    "correlations",  
    "mean.diffs"  
  ),  
  abs = FALSE,  
  var.order = "unadjusted",
```

```

thresholds = c(cor = 0.1, m = 0.1),
line = TRUE,
)

```

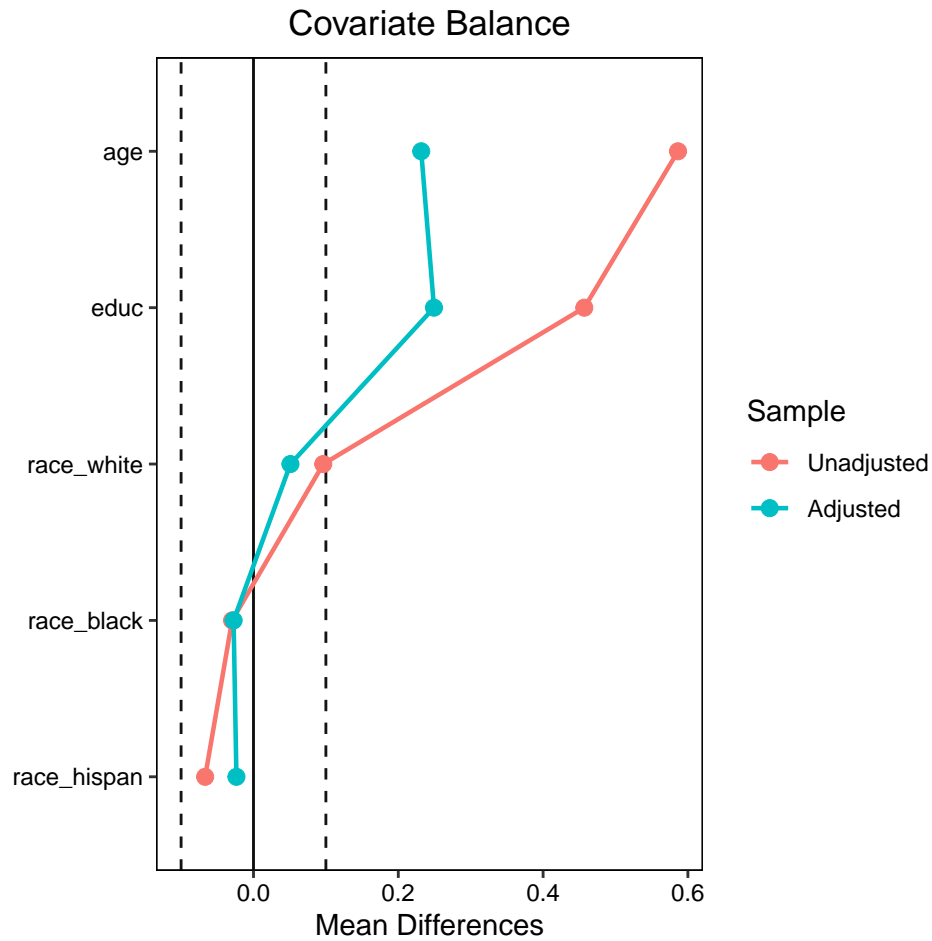


Figure 2: Love plot

Once we are satisfied with the estimated balancing weights, we can proceed to fit our (weighted) outcome model.

```

fit <- glm(
  formula = re78 ~ treat + age + educ + nodegree + race,
  weights = est_sel_weights_trim$weights,
  data = dat
)

```

)

2.2 The SelectionWeights R package

TBD.

3 Suggestions

We would like to further provide some suggestions based on previous experience.

- It is useful to compare the results (e.g., predicted outcome or exposure coefficient) with and without IP weights, to check whether something went wrong during the estimation process. This information can be put in the supplementary material.
- It might be useful to provide further information about the balancing. More tools are available in the `cobalt` R package (Greifer 2024a).
- Use robust SEs.
- By cohort estimation.

References

- Alten, Sjoerd van, Benjamin W Domingue, Titus Galama, and Andries T Marees. 2022. “Reweighting the UK Biobank to Reflect Its Underlying Sampling Population Substantially Reduces Pervasive Selection Bias Due to Volunteering.” *medRxiv*, 2022–05.
- Greifer, Noah. 2024a. *Cobalt: Covariate Balance Tables and Plots*. <https://ngreifer.github.io/cobalt/>.
- . 2024b. *WeightIt: Weighting for Covariate Balance in Observational Studies*. <https://ngreifer.github.io/WeightIt/>.
- Hernán, Miguel A, and James M Robins. 2010. “Causal Inference.” CRC Boca Raton, FL.